



# How to build a data pipeline without data

*Synthetic data generation  
and testing with Python*

# About me

Hi there, I'm Ruan Pretorius 🙋

- ☕ I turn coffee into data pipelines and AI
- 💻 I am a data scientist at [melio.ai](https://melio.ai)
  - We help you build and deploy your data intensive apps to unlock value from your data
- 🔗 You can find me on GitHub [@ruankie](https://github.com/ruankie)
- ✉ Or contact me via email: [ruan@melio.ai](mailto:ruan@melio.ai)



## Outline

- What are data pipelines and why do we need them?
- Challenges of building and testing data pipelines
- How to use synthetic data to test data pipelines
- Tools and methods to use when generating reliable synthetic data in Python
- Benefits and challenges of using synthetic data for testing data pipelines

## What is a data pipeline?

- A data pipeline is a series of operations used to extract, load, transform, validate, or write data
- From various sources into a target file system, database, or data warehouse

## Data pipelines without real data

- Sometimes, we may not have access to the real data that we want to process in our data pipeline.
- It could be:
  - Sensitive or confidential and can't be shared
  - Not yet collected or available
  - Too large or complex to handle for initial testing

## Data pipelines without real data

- Without real data, it is challenging to:
  - Design and build downstream apps that consume the data
  - Develop the data extract, transform, and load (ETL) logic
  - Test the functionality and performance of the data pipeline

## Synthetic data to test data pipelines

- Synthetic data is artificially generated data that mimics the characteristics and behavior of real data
- Synthetic data can help us to test our data pipelines by:
  - Providing realistic sample data
  - Allowing control of the size, shape, and distribution of the data
  - Enabling simulations of different scenarios and edge cases
  - Reducing the risk of exposing sensitive or confidential information

# Demo

- In this demo, I'll show you how you can create synthetic data
- Using a Python package called `Faker`
- And how to use `Flyway` to load the synthetic data into a `Postgres` database for repeatable deployments
- So that you can test your pipelines without real data





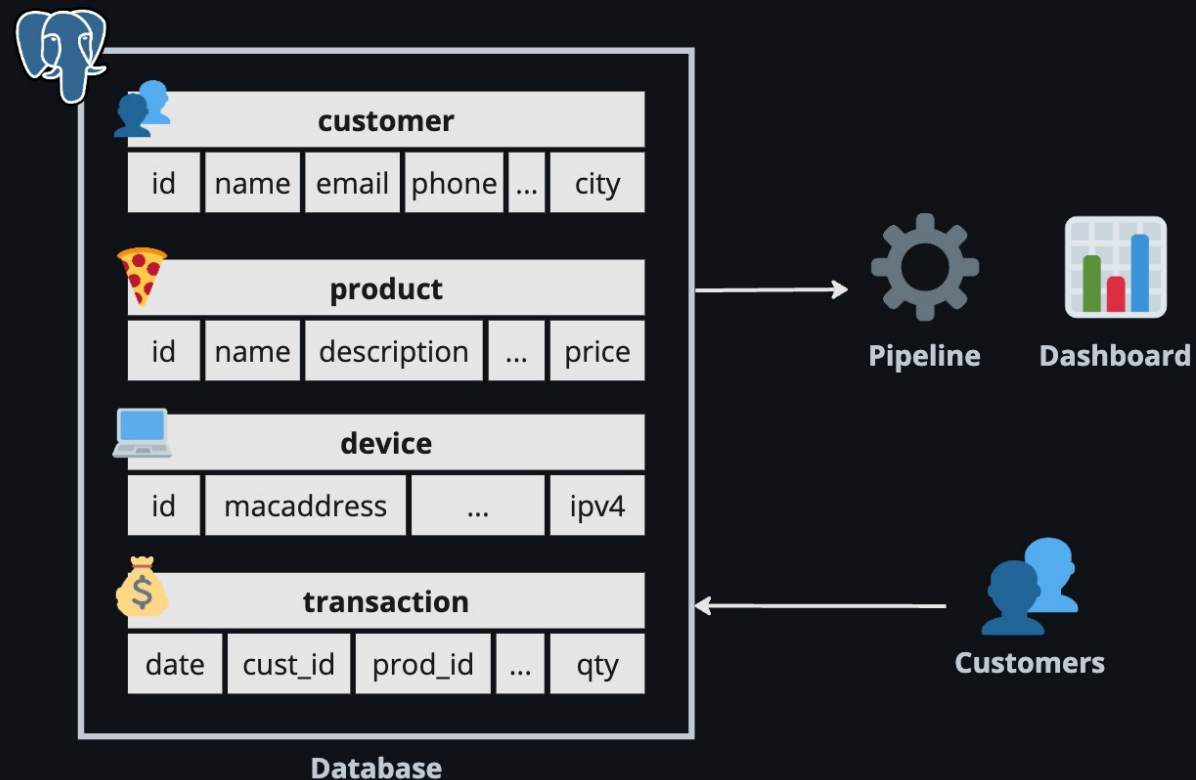


## Our scenario

- Let's pretend we just started a new e-commerce website
- We have an idea of what kind of data we'll have for
  - Customers
  - Products
  - Transactions

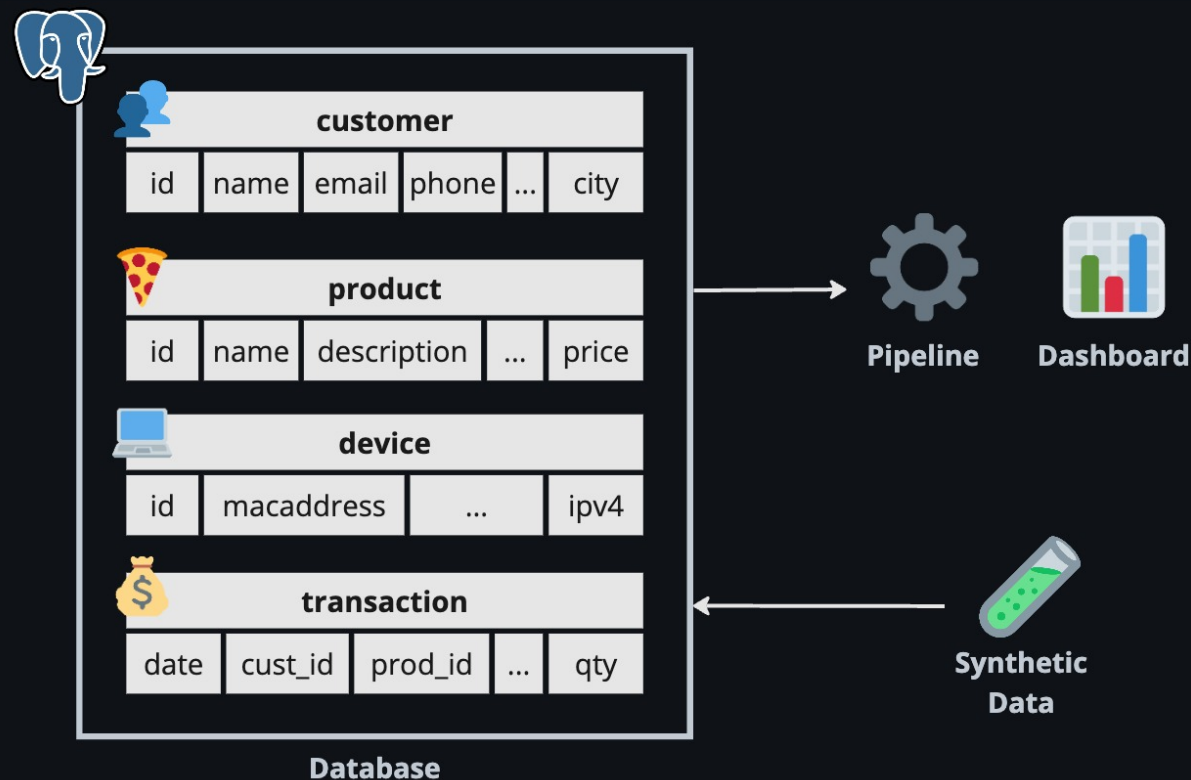
## The data problem

- Now we want to start building different data pipelines and visualisations to see how well our business is doing
- We want our systems to work as soon as we get customers



## The data problem

- But we don't have customers yet (or data for them)
- So let's make some
- Then we can build everything downstream and it should work when we get real data



## Install and import tools

```
pip install SQLAlchemy Faker
```

- **SQLAlchemy** to create database objects

```
from sqlalchemy import Column, Integer, String, DateTime
from sqlalchemy.orm import declarative_base
Base = declarative_base()
```

- **Faker** to generate synthetic data

```
from faker import Faker
fake = Faker()
```

## Customer object

Class to store customer information

```
class Customer(Base):  
    __tablename__ = "customers"  
    id = Column(Integer, primary_key=True)  
    name = Column(String(100))  
    email = Column(String(100))  
    phone = Column(String(25))  
    address = Column(String(250))  
    city = Column(String(100))  
    country = Column(String(100))
```

## Customer data

Customer generator using **Faker** for synthetic data

```
def generate_customer(id: int):  
    customer = Customer(  
        id=id,  
        name=fake.name(),  
        email=fake.email(),  
        phone=fake.phone_number(),  
        address=fake.street_address(),  
        city=fake.city(),  
        country=fake.country()  
    )  
    return customer
```

## Product object

Class to store product information

```
class Product(Base):  
    __tablename__ = "products"  
    id = Column(Integer, primary_key=True)  
    name = Column(String(50))  
    description = Column(String(200))  
    category = Column(String(50))  
    price = Column(Numeric(10, 2))
```

## Product data

Product generator using **Faker** for synthetic data

```
def generate_product(id: int):  
    product = Product(  
        id=id,  
        name=fake.word(),  
        description=fake.sentence(),  
        category=fake.random_element(  
            elements=("Electronics", "Fashion", "Books", "Games", "Sports", "Food")  
        ),  
        price=fake.pydecimal(left_digits=3, right_digits=2, positive=True)  
    )  
    return product
```



## Device object

Class to store device information

```
class Device(Base):  
    __tablename__ = "devices"  
    id = Column(Integer, primary_key=True)  
    platform = Column(String(250))  
    ipv4 = Column(String(50))  
    macaddress = Column(String(50))
```

## Device data

Device generator using **Faker** for synthetic data

```
def generate_device(id: int):  
    device = Device(  
        id=id,  
        platform=fake.user_agent(),  
        ipv4=fake.ipv4(),  
        macaddress=fake.mac_address()  
    )  
    return device
```

# Transaction object

Class to store transaction information

```
class Transaction(Base):  
    __tablename__ = "transactions"  
    id = Column(Integer, primary_key=True)  
    date_time = Column(DateTime)  
    customer_id = Column(Integer)  
    product_id = Column(Integer)  
    quantity = Column(Integer)  
    device_id = Column(Integer)  
    payment_method = Column(String(50))
```

# Transaction data

Transaction generator using **Faker** for synthetic data

```
def generate_transaction(
    id: int,
    customers: list[Customer],
    products: list[Product],
    devices: list[Device]
):
    tr = Transaction(
        id=id,
        date_time=fake.date_between(start_date=START_DATE, end_date=END_DATE),
        customer_id=random.choice(customers).id,
        product_id=random.choice(products).id,
        quantity=fake.random_int(min=1, max=20),
        device_id=random.choice(devices).id,
        payment_method=fake.random_element(
            elements=("Credit Card", "EFT", "Bitcoin", "Reward Points")
        )
    )
    return tr
```

## Generate data

Use our functions to generate synthetic data

```
customers = [generate_customer(i) for i in range(1000)]
products = [generate_product(i) for i in range(60)]
devices = [generate_device(i) for i in range(1000)]

transactions = [
    generate_transaction(i, customers, products, devices) for i in range(5000)
]
```

# Write data do database

We have a choice

- Just use `SQLAlchemy` to write to our database
- Or use `Flyway` (we'll use this option)
  - It handles version control of our SQL scripts
  - And handles database migrations
    - For repeatable deployments
    - And certainty about our database state
  - This all makes it easier to collaborate with other developers

## Generate SQL scripts

- `Flyway` is almost like `git` for your database
- So we need to get the `CREATE` and `INSERT` statements
  - That will create our tables
  - And to insert our synthetic data

# Generate SQL scripts

Create SQL string > Write to `.sql` file

- You can do some simple string manipulation:

```
sql = f"""CREATE TABLE {table} (  
    id SERIAL NOT NULL,  
    name VARCHAR(100) NOT NULL  
);"""
```

- Or you can use `SQLAlchemy`:

```
from sqlalchemy import create_engine  
from sqlalchemy.schema import CreateTable  
  
engine = create_engine("postgresql:///memory:")  
sql = str(CreateTable(cls.__table__).compile(engine))
```



## Generate SQL scripts

Create SQL string > Write to `.sql` file

- You can do some simple string manipulation:

```
sql = f"INSERT INTO {table} VALUES ({object.id}, {object.name})"
```

- Or you can use `SQLAlchemy`:

```
from sqlalchemy.sql.expression import insert

insert_stmt = insert(cls.__table__)\
    .values(records)\ # list of records from helper function
    .compile(compile_kwargs={"literal_binds": True})
sql = str(insert_stmt)
```

# 🦋 Use Flyway to create tables and data

Install the CLI ([instructions online](#))

- Configure Flyway by creating a new project and specifying database connections
- Drop all tables, views, procedures etc. in the configured schemas to start fresh

```
flyway clean
```

- Migrate schemas to the latest version by applying SQL scripts

```
flyway migrate
```

## 🦋 Use Flyway to create tables and data

After downloading and installing Flyway Desktop

- Configure Flyway by creating a new project and specifying database connections

# 🦋 Use Flyway to create tables and data

To start fresh, perform a `clean` to drop all tables, views, procedures etc. in the configured schemas

The screenshot shows the Flyway Desktop application window. The title bar reads "Flyway Desktop". The top navigation bar includes the Flyway logo, "Community", a tab for "pyconza23-demo", and buttons for "Feedback", "Preview", settings, user profile, help, and notifications. The "redgate" logo is in the top right corner.

The main area is titled "Migrations" and features a search bar, a checkbox for "Only show pending migrations", and refresh and add buttons. Below this is a table of migrations:

Category	Version	Description	Type	Undoable	Date migrated	State	Execution time
Versioned	1.1	create customers table	SQL	✗	2023-09-24 15:53:21	✓ Success	39ms
Versioned	1.2	insert customers records	SQL	✗	2023-09-24 15:53:21	✓ Success	39ms
Versioned	1.3	create products table	SQL	✗	2023-09-24 15:53:21	✓ Success	20ms
Versioned	1.4	insert products records	SQL	✗	2023-09-24 15:53:21	✓ Success	11ms
Versioned	1.5	create devices table	SQL	✗	2023-09-24 15:53:21	✓ Success	18ms

The right sidebar contains a database selector showing "default" and "public" schemas, with a link to "Configure target database...". Below this, it states "0 pending migrations" and offers to "View command (0 parameters)". There is a "Migrate" dropdown and a "Run migrate" button. At the bottom, there is an "Advanced settings" section with a link to "View Flyway configuration settings".

# 🦋 Use Flyway to create tables and data

Perform a `migrate` to update schemas to the latest version by applying SQL scripts

The screenshot shows the Flyway Desktop application window. The title bar reads "Flyway Desktop". The top navigation bar includes the Flyway logo, "Community", a tab for "pyconza23-demo", and buttons for "Feedback", "Preview", settings, user profile, help, and notifications, along with the "redgate" logo.

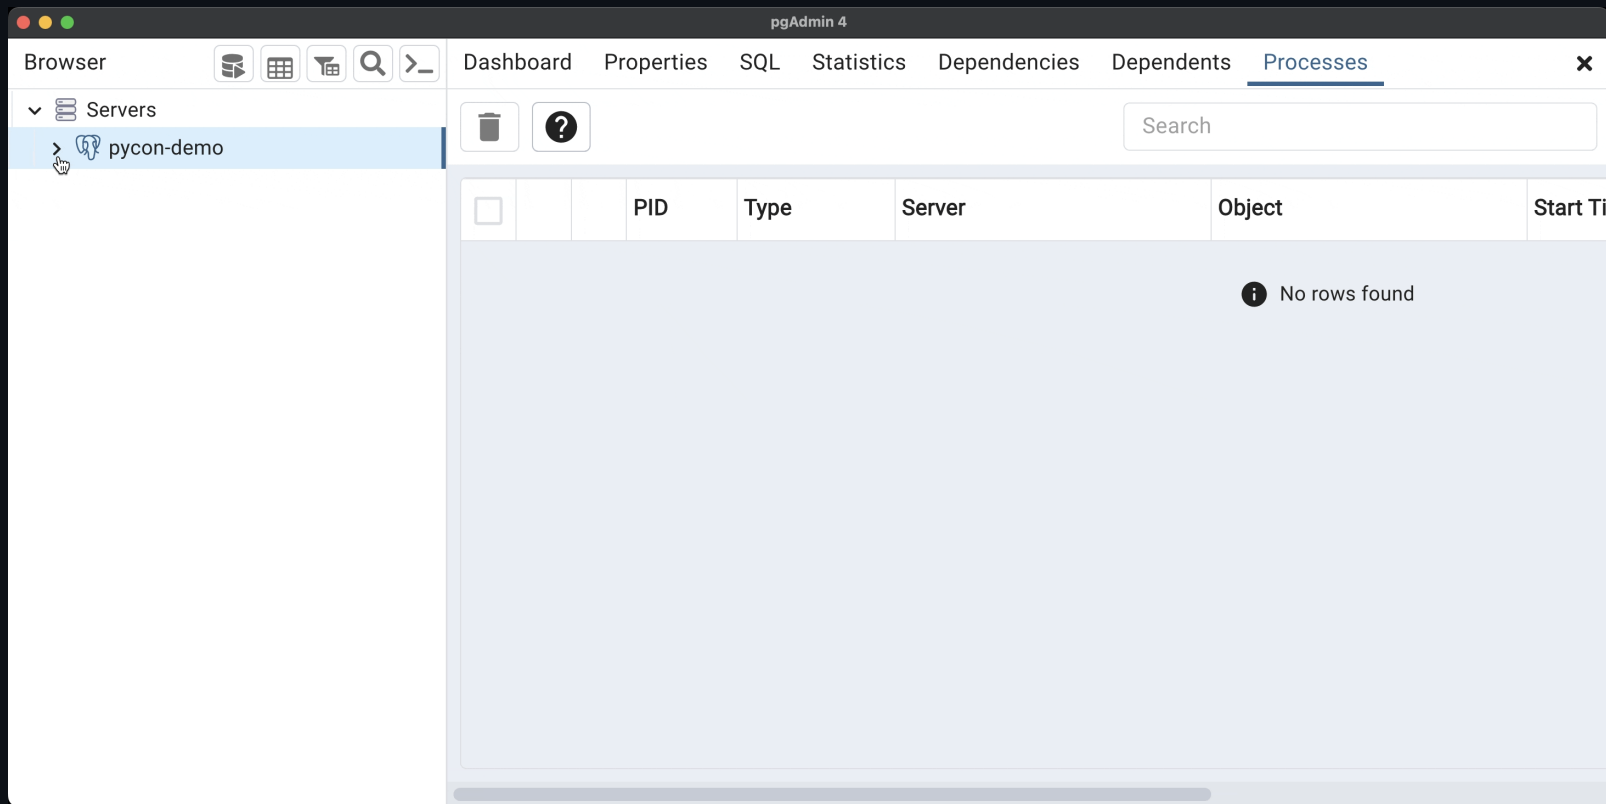
The main area is titled "Migrations" and contains a search bar, a checkbox for "Only show pending migrations", and a refresh button. Below this is a table of migrations:

Category	Version	Description	Type	Undoable	Date migrated	State	Execution time
Versioned	1.1	create customers table	SQL	✗	-	Pending	
Versioned	1.2	insert customers records	SQL	✗	-	Pending	
Versioned	1.3	create products table	SQL	✗	-	Pending	
Versioned	1.4	insert products records	SQL	✗	-	Pending	
Versioned	1.5	create devices table	SQL	✗	-	Pending	

The right sidebar shows the database configuration with a dropdown set to "default" and "public". It indicates "8 pending migrations" and provides options to "View command (0 parameters)", "Clean", "Run clean", and "Advanced settings".

# 👁 Inspect data

Inspect database with `pgAdmin` to see if your synthetic data is ready



## Pros and cons of synthetic data

- Benefits:
  - It can speed up the development and testing process
  - It can increase the coverage and quality of testing
- Challenges & limitations:
  - It may not capture all the nuances and variations of real data
  - It may require additional effort and resources to create and maintain synthetic data

## Tips & best practices

- Define the scope and purpose of your synthetic data before starting
- Be careful of wasting time trying to make perfectly realistic data
- Use existing tools and libraries to generate synthetic data where possible
- Validate and verify your synthetic data against your real data schema and business rules
- Document your synthetic data generation process and code






## Summary

- We discussed some challenges of building and testing data pipelines without real data
- We learned how to build synthetic data in Python to test our data pipelines
- We also showed how we used Flyway to load the synthetic data into a Postgres database



# Thank you!

-  GitHub: [@ruankie](#)
-  Email: [ruan@melio.ai](mailto:ruan@melio.ai)
-  Melio website: [melio.ai](https://melio.ai)

